

# BAJANDO EL NIVEL

Metiendo mano en el kernel Linux.

samuel (4nt1) paschuan



Secure  
Jump



# ¿quien soy?

cofounder @ securejump

pentester @ pentest spa, sombreroblanco

director @ CUC 2025

“director” @ RESACA Chile

profesor/ayudante @ umayor

hoy hablaremos de...

compilación.

depuración.

tracing.

una intro antes

intro.

# syscalls

## interacción user-kernel

kernelspace  
donde vive el kernel

userspace  
donde vive el usuario

locking  
tipo de mutex

mutex  
objeto de sincronización

rcu  
read, copy, update

irq  
interrupt request

shed  
planificador de ejecución

hardware rings  
seguridad física

VSS  
virtual subsystem

MMU  
memory management unit

compilación.

¿cómo se compila el kernel?

así.

```
make defconfig  
make -j$(nproc)
```

no es suficiente.

(para nuestro use-case)

flags de depuración,  
tracing y tooling.

lo mínimo sería...

```
make defconfig
# o
make kvm_guest.config
# .config
CONFIG_DEBUG_INFO_DWARF_TOOLCHAIN_DEFAULT=y
CONFIG_KGDB=y
CONFIG_FTRACE=y
CONFIG_RANDOMIZE_BASE=n
make olddefconfig
make -j$(nproc)
```

funcional... si.

pero subóptimo.

entonces . . .

¡ ¿cómo? !

```
make defconfig
# o
make kvm_guest.config
# vim .config
CONFIG_DEBUG_INFO_DWARF_TOOLCHAIN_DEFAULT=y
CONFIG_KGDB=y
CONFIG_FTRACE=y
CONFIG_BOOTTIME_TRACING=y
CONFIG_FUNCTION_TRACER=y
CONFIG_STACK_TRACER=y
CONFIG_FTRACE_SYSCALLS=y
CONFIG_SECURITYFS=y
CONFIG_RANDOMIZE_BASE=n
CONFIG_CONFIGFS_FS=y
CONFIG_STRICT_KERNEL_RWX=n
CONFIG_KGDB_SERIAL_CONSOLE=y
CONFIG_GDB_SCRIPTS=y
make olddefconfig
```

este setup es mejor . . .

sin embargo...

solamente tenemos depuración . . .

y tracing...

esto es por diseño.

el tooling dependerá...

del caso a caso.

¿el repertorio?

el siguiente.

# KASAN. Kernel Address Sanitizer

# KFENCE. Kernel Electronic Fence

# KPROBE. Kernel Probes

# KCOV.

# Kernel Coverage

# KMSAN. Kernel Memory Sanitizer

# UBSAN. Undefined Behaviour Sanitizer

# KMLD.

# Kernel Memory Leak Detector

# KCSAN. Kernel Concurrency Sanitizer

# KUNIT. Kernel Unit Testing

# KTAP. Kernel Test Anything Protocol

# UAPIC. UAPI Checker

# KGSLA.

# Kernel GPIO Sloppy Logic Analyzer

dado el tiempo...

solo veremos una.

# KASAN. Kernel Address Sanitizer

habilitemosla.

```
# vim .config
CONFIG_KASAN=y
CONFIG_KASAN_GENERIC=y
make olddefconfig
make -j$(nproc)
# recommended cmdline: kasan.fault=panic
```

KASAN sirve para dos cosas

detección de UAFs...

y 00Bs

lo veremos más adelante.

depuración.

tenemos un kernel.

pero necesitamos usarlo.

pero... ¿cómo?

qemu-system-x86\_64

la forma básica:

```
qemu-system-x86_64 \
-kernel arch/x86_64/boot/bzImage \
-m 2G \
-s \
-S
```

ahora , abramos gdb .

```
gdb -ex 'target remote :1234' vmlinux  
continue
```

lo que logramos no es mucho.

```
[ 3.155688] ? kmem_cache_alloc_noprof+0x130/0x300
[ 3.155737] ? __asan_memcpy+0x3c/0x60
[ 3.155776] ? getname_kernel+0x148/0x350
[ 3.155822] mount_root+0x29d/0x2f0
[ 3.155862] ? __kasan_slab_alloc+0x63/0x80
[ 3.155904] ? __px__mount_root+0x10/0x10
[ 3.155948] ? __px__init_mknod+0x10/0x10
[ 3.155994] ? __asan_memcpy+0x3c/0x60
[ 3.156042] ? getname_kernel+0x148/0x350
[ 3.156085] prepare_namespace+0x216/0x280
[ 3.156133] ? __px__prepare_namespace+0x10/0x10
[ 3.156180] kernel_init_freeable+0x4b5/0x4d0
[ 3.156227] ? __px__kernel_init+0x10/0x10
[ 3.156275] kernel_init+0x19/0x150
[ 3.156313] ? __px__kernel_init+0x10/0x10
[ 3.156353] ? __px__kernel_init+0x10/0x10
[ 3.156395] ret_from_fork+0x10d/0x1b0
[ 3.156436] ? __px__kernel_init+0x10/0x10
[ 3.156478] ret_from_fork_asm+0x1a/0x30
[ 3.156526] </TASK>
[ 3.156968] Kernel Offset: 0x33800000 from 0xffffffff81000000 (relocation range: 0xffffffff80000000-0xffffffffbfffffff)
[ 3.157414] ---[ end Kernel panic - not syncing: VFS: Unable to mount root fs on unknown-block(0,0) ]---
```

tenemos varias opciones.

generar nuestra propia CPI0...

o una de syzkaller.

veamos cómo generar una CPIO.

```
# NECESITAMOS!!! un binario ELF o script llamado "init".
# en mi caso, compilé Bash estáticamente para esta tarea.
wget https://ftp.gnu.org/gnu/bash/bash-5.3.tar.gz
tar xfv bash-5.3.tar.gz
cd bash5.3
./configure --enable-static-link
# para la build estática, se requieren librerías GCC estáticas.
# busca en Google cuales son las que te sirven en tu distro.
make clean
make -j$(nproc)
mv bash init
mkdir cpio
mv init cpio
cd cpio
ls | cpio -H newc -o > initram.cpio
cd linux/source/code/path
qemu-system-x86_64 \
    -m 2G \
    -kernel arch/x86_64/boot/bzImage \
    -initrd /path/to/initram.cpio \
    -s \
    -S
```

```
[ 3.539696] Loaded X.509 cert 'sforshee: 00b28ddf47aef9cea?'
[ 3.541237] Loaded X.509 cert 'wens: 61c038651aabdcf94bd0ac7ff06c7248db18c600
',
[ 3.543285] ALSA device list:
[ 3.543393]   No soundcards found.
[ 3.547553] faux_driver regulatory: Direct firmware load for regulatory.db fa
iled with error -2
[ 3.548103] cfg80211: failed to load regulatory.db
[ 3.612468] Freeing unused kernel image (initmem) memory: 5904K
[ 3.615687] input: ImExPS/2 Generic Explorer Mouse as /devices/platform/i8042
/serio1/input/input3
[ 3.618351] Write protecting the kernel read-only data: 67584k
[ 3.620377] Freeing unused kernel image (text/rodata gap) memory: 644K
[ 3.621154] Freeing unused kernel image (rodata/data gap) memory: 600K
[ 3.740652] x86/mm: Checked W+X mappings: passed, no W+X pages found.
[ 3.741131] Run /init as init process
init: cannot set terminal process group (-1): Inappropriate ioctl for device
init: no job control in this shell
init-5.3# echo hola desde el kernel    8.870364] random: crng init done

hola desde el kerne
init-5.3# echo hola desde el kernel
init-5.3# echo hola desde el kernel
hola desde el kernel
init-5.3#
```

mientras esté compilado  
estáticamente...

todo puede ser un init.

flags útiles de qemu.

-nographic  
todo a través de la terminal

-append  
añade parámetros a la cmdline

-enable-kvm  
usa KVM en vez de TCG.

-pidfile file.pid  
almacena el proceso en un PIDfile

-monitor  
envio del monitor hacia un dev

-net  
configura un dispositivo de red

(s)breaks have no(t much)  
power here

muchos usamos el querido break

pero . . .

sin MMU, vmem . . .

no tiene poder alguno.

```
[ 3.651404] cfg80211: Loading compiled-in X.509 certificates for regulatory database
[ 3.662009] kworker/u4:1 (50) used greatest stack depth: 29360 bytes left
[ 3.665892] kworker/u4:0 (51) used greatest stack depth: 29024 bytes left
[ 3.676660] Loaded X.509 cert 'sforshee: 00b28ddf47aef9cea?'
[ 3.678398] Loaded X.509 cert 'wens: 61c038651aabdcf94bd0ac7ff06c7248db18c600
,
[ 3.681049] ALSA device list:
[ 3.681273]   No soundcards found.
[ 3.683837] faux_driver regulatory: Direct firmware load for regulatory.db failed with error -2
[ 3.684167] cfg80211: failed to load regulatory.db
[ 3.749524] Freeing unused kernel image (initmem) memory: 5904K
[ 3.752565] input: ImExPS/2 Generic Explorer Mouse as /devices/platform/i8042/serio1/input/input3
[ 3.754902] Write protecting the kernel read-only data: 67584K
[ 3.756899] Freeing unused kernel image (text/rodata gap) memory: 644K
[ 3.757633] Freeing unused kernel image (rodata/data gap) memory: 600K
[ 3.894227] x86/mm: Checked W+X mappings: passed, no W+X pages found.
[ 3.894677] Run /init as init process
/init: cannot set terminal process group (-1): Inappropriate ioctl for device
/init: no job control in this shell
init-5.3# echo y mi break?
y mi break?
init-5.3# _
```

```
(gdb) break kernel_init
```

```
Breakpoint 1 at 0xffffffff8453a000: file init/main.c, line 1473.
```

```
(gdb) c
```

```
Continuing.
```

su funcionamiento es irregular

a veces funciona . . .

```
SeaBIOS (version 1.17.0-5.fc42)
```

```
iPXE (https://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+7EFCAE60+7EF0AE60 CA00
```

```
Booting from ROM...
No EFI environment detected.
early console in extract_kernel
input_data: 0x0000000004e6b2c4
input_len: 0x0000000001c25a9f
output: 0x0000000001000000
output_len: 0x0000000005a15c90
kernel_total_size: 0x0000000005a28000
needed_size: 0x0000000005c00000
trampoline_32bit: 0x0000000000000000
```

```
Decompressing Linux... Parsing ELF... done.
Booting the kernel (entry_offset: 0x000000000355e2f0).
```

```
(gdb) b start_kernel
Breakpoint 1 at 0xffffffff86244430: file init/main.c, line 910.
(gdb) c
Continuing.
```

```
Breakpoint 1, start_kernel () at init/main.c:910
910
(gdb) █
```

a veces no.

```
[ anti@mothership ]  
01:33 PM  
~/Documents/Research/linux-debugging/linux-6.17.7 $ cat ../qemu.sh | bash
```

---

```
(gdb) break start_kernel  
Breakpoint 1 at 0xffffffff86244430: file init/main.c, line 910.  
(gdb) c  
Continuing.  
Warning:  
Cannot insert breakpoint 1.  
Cannot access memory at address 0xffffffff86244430  
  
Command aborted.  
(gdb) █
```

y cuando no . . .

usamos (h)breaks.

hardware assisted breaks.

```
No EFI environment detected.  
early console in extract_kernel  
input_data: 0x0000000004e6b2c4  
input_len: 0x0000000001c25a9f  
output: 0x0000000001000000  
output_len: 0x0000000005a15c90  
kernel_total_size: 0x0000000005a28000  
needed_size: 0x0000000005c00000  
trampoline_32bit: 0x0000000000000000
```

```
Decompressing Linux... Parsing ELF... done.  
Booting the kernel (entry_offset: 0x00000000355e2f0).
```

---

```
(gdb) hb start_kernel  
Note: breakpoint 1 (disabled) also set at pc 0xffffffff86244430.  
Hardware assisted breakpoint 2 at 0xffffffff86244430: file init/main.c, line 910.  
(gdb) c  
Continuing.
```

  

```
Breakpoint 2, start_kernel () at init/main.c:910  
910 {  
(gdb) █
```

anatomía de un driver.

los drivers de Linux...

a alto nivel...

son sencillos.

veaamos el esqueleto de uno.

```
#include<linux/module.h>
#include<linux/init.h>

int kmod_init(void)
{
    printk("kmod-hello - module loaded!\n");
    return 0;
}

void kmod_exit(void)
{
    printk("kmod-hello - module unloaded!\n");
}

module_init(kmod_init);
module_exit(kmod_exit);

MODULE_LICENSE("GPL");
```

si bien solo el esqueleto...

esta estructura es global

función `module_init` con un  
valor `int` de retorno

función `module_exit` sin valor  
de retorno

lo que si es necesario...

es el MODULE\_LICENSE.

esto permite impedir la carga  
de modulos propietarios.

si... impedir.

major-minors.

los major-minor values...

son importantes.

cada driver tiene un major.

cada dispositivo un minor.

generalmente asignados  
dinámicamente.

algunos estáticamente.

```
~ $ ls -l /dev/tty*
crw-rw-rw-. 1 root tty      5,  0 Nov 17 16:31 /dev/tty
crw--w----. 1 root tty      4,  0 Nov 14 20:25 /dev/tty0
crw--w----. 1 root tty      4,  1 Nov 14 20:25 /dev/tty1
crw--w----. 1 root tty      4, 10 Nov 14 20:25 /dev/tty10
crw--w----. 1 root tty      4, 11 Nov 14 20:25 /dev/tty11
[SNIP]
```

estructuras driverescas.

**file\_operations**  
from: include/linux/fs.h

```
2156 struct file_operations {
2157     struct module *owner;
2158     fop_flags_t fop_flags;
2159     loff_t (*llseek) (struct file *, loff_t, int);
2160     ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
2161     ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
2162     ssize_t (*read_iter) (struct kiocb *, struct iov_iter *);
2163     ssize_t (*write_iter) (struct kiocb *, struct iov_iter *);
2164     int (*iopoll)(struct kiocb *kiocb, struct io_comp_batch *,
2165                   unsigned int flags);
2166     int (*iterate_shared) (struct file *, struct dir_context *);
2167     __poll_t (*poll) (struct file *, struct poll_table_struct *);
2168     long (*unlockd_ioctl) (struct file *, unsigned int, unsigned long);
2169     long (*compat_ioctl) (struct file *, unsigned int, unsigned long);
2170     int (*mmap) (struct file *, struct vm_area_struct *);
2171     int (*open) (struct inode *, struct file *);
2172     int (*flush) (struct file *, fl_owner_t id);
2173     int (*release) (struct inode *, struct file *);
2174     int (*fsync) (struct file *, loff_t, loff_t, int datasync);
2175     int (*fasync) (int, struct file *, int);
2176     int (*lock) (struct file *, int, struct file_lock *);
2177     unsigned long (*get_unmapped_area)(struct file *, unsigned long, unsigned long, unsigned long, unsigned long);
2178     int (*check_flags)(int);
2179     int (*flock) (struct file *, int, struct file_lock *);
2180     ssize_t (*splice_write)(struct pipe_inode_info *, struct file *, loff_t *, size_t, unsigned int);
2181     ssize_t (*splice_read)(struct file *, loff_t *, struct pipe_inode_info *, size_t, unsigned int);
2182     void (*splice_eof)(struct file *file);
2183     int (*setlease)(struct file *, int, struct file_lease **, void **);
2184     long (*fallocate)(struct file *file, int mode, loff_t offset,
2185                       loff_t len);
2186     void (*show_fdinfo)(struct seq_file *m, struct file *f);
2187 #ifndef CONFIG_MMU
2188     unsigned (*mmap_capabilities)(struct file *);
2189 #endif
2190     ssize_t (*copy_file_range)(struct file *, loff_t, struct file *,
2191                               loff_t, size_t, unsigned int);
2192     loff_t (*remap_file_range)(struct file *file_in, loff_t pos_in,
2193                               struct file *file_out, loff_t pos_out,
2194                               loff_t len, unsigned int remap_flags);
2195     int (*fadvise)(struct file *, loff_t, loff_t, int);
2196     int (*uring_cmd)(struct io_uring_cmd *ioucmd, unsigned int issue_flags);
2197     int (*uring_cmd_iopoll)(struct io_uring_cmd *, struct io_comp_batch *,
2198                           unsigned int poll_flags);
2199     int (*mmap_prepare)(struct vm_area_desc *);
2200 } __randomize_layout;
```

```
static struct file_operations my_null_fops = {
    .owner = THIS_MODULE,
    .read = my_null_read,
    .write = my_null_write,
    .open = my_null_open,
    .release = my_null_release,
};
```

file  
from: include/linux/fs.h

```
1100 struct file {
1101         spinlock_t          f_lock;
1102         fmode_t             f_mode;
1103         const struct file_operations *f_op;
1104         struct address_space *f_mapping;
1105         void                *f_private_data;
1106         struct inode         *f_inode;
1107         unsigned int         f_flags;
1108         unsigned int         f_iocb_flags;
1109         const struct cred    *f_cred;
1110         struct fown_struct  *f_owner;
1111         /* --- cacheline 1 boundary (64 bytes) --- */
1112         struct path           f_path;
1113         union {
1114             /* regular files (with FMODE_ATOMIC_POS) and directories */
1115             struct mutex        f_pos_lock;
1116             /* pipes */
1117             u64                 f_pipe;
1118         };
1119         loff_t               f_pos;
1120 #ifdef CONFIG_SECURITY
1121         void                *f_security;
1122 #endif
1123         /* --- cacheline 2 boundary (128 bytes) --- */
1124         errseq_t              f_wb_err;
1125         errseq_t              f_sb_err;
1126 #ifdef CONFIG_EPOLL
1127         struct hlist_head     *f_ep;
1128 #endif
1129         union {
1130             struct callback_head f_task_work;
1131             struct llist_node   f_llist;
1132             struct file_ra_state f_ra;
1133             freeptr_t           f_freeptr;
1134         };
1135         file_ref_t            f_ref;
1136         /* --- cacheline 3 boundary (192 bytes) --- */
1137 } __randomize_layout
1138 __attribute__((aligned(4))); /* lest something weird decides that 2 is OK */
1139
```

`inode`  
from: `include/linux/fs.h`

```
struct inode {  
    [SNIP]  
    dev_t i_rdev;  
    struct cdev *i_cdev;  
    [SNIP]  
}
```

para entender esto...

debemos entender . . .

la VFS! Virtual File System

la cual no veremos hoy.

¡pero!

pregunta de investigación

cada fs tiene su propio driver

EXT4, BTRFS, etc.

sus propias funciones.

`ext4_file_open,`  
`btrfs_file_open, etc.`

pero nosotros no usamos esas  
funciones.

¿cómo?

¿por qué?

¿y qué tiene que ver la VFS?

depurando un use-after-free.

veamos el siguiente código:

```
#include<linux/module.h>
#include<linux/init.h>
#include<linux/slab.h>

int use_after_free(void)
{
    printk("kmod-uaf - module loading!\n");
    char *leak = kzalloc(100, GFP_KERNEL);
    if (!leak)
    {
        printk(KERN_ERR "Memory allocation failed\n");
        goto out;
    }
    strcpy(leak, "This is a use-after-free example");
    kfree(leak);
    printk(KERN_INFO "Using freed memory: %s\n", leak);
out:
    return 0;
}

void kmod_exit(void)
{
    printk("kmod-uaf - module unloaded!\n");
}

module_init(use_after_free);
module_exit(kmod_exit);

MODULE_LICENSE("GPL");
```

```
#include<linux/module.h>
#include<linux/init.h>
#include<linux/slab.h>

int use_after_free(void)
{
    printk("kmod-uaf - module loading!\n");
    char *leak = kzalloc(100, GFP_KERNEL);
    if (!leak)
    {
        printk(KERN_ERR "Memory allocation failed\n");
        goto out;
    }
    strcpy(leak, "This is a use-after-free example");
    kfree(leak);
    printk(KERN_INFO "Using freed memory: %s\n", leak);
out:
    return 0;
}

void kmod_exit(void)
{
    printk("kmod-uaf - module unloaded!\n");
}

module_init(use_after_free);
module_exit(kmod_exit);

MODULE_LICENSE("GPL");
```

el uaf, en este caso, es  
aparente.

pero no siempre es así.

desarrollando . . .

¿cómo lo encontramos?

con KASAN.

y con la cmdline  
**kasan.fault=panic**.

```
[ 0.768208] scsi 1:0:0:0: CD-ROM QEMU QEMU DVD-ROM 2.5+ PQ: 0 ANSI: 5
[ 0.769487] sd 0:0:0:0: [sda] Preferred minimum I/O size 512 bytes
[ 0.781207] sd 0:0:0:0: [sda] Attached SCSI disk
[ 0.795243] sr 1:0:0:0: [sr0] scsi3-mmc drive: 4x/4x cd/rw xa/form2 tray
[ 0.796341] cdrom: Uniform CD-ROM driver Revision: 3.20
[ 0.814677] sr 1:0:0:0: Attached scsi generic sg1 type 5
[ 0.946200] e1000 0000:00:03.0 eth0: (PCI:33MHz:32-bit) 52:54:00:12:34:56
[ 0.947302] e1000 0000:00:03.0 eth0: Intel(R) PRO/1000 Network Connection
[ 0.948390] e1000e: Intel(R) PRO/1000 Network Driver
[ 0.949128] e1000e: Copyright(c) 1999 - 2015 Intel Corporation.
[ 0.950056] sky2: driver version 1.30
[ 0.950981] usbcore: registered new interface driver usblp
[ 0.951846] usbcore: registered new interface driver usb-storage
[ 0.952800] i8042: PNP: PS/2 Controller [PNP0303:KBD,PNP0f13:MOU] at 0x60,0x64 irq 1,12
[ 0.954892] serio: i8042 KBD port at 0x60,0x64 irq 1
[ 0.955658] serio: i8042 AUX port at 0x60,0x64 irq 12
[ 0.957136] input: AT Translated Set 2 keyboard as /devices/platform/i8042/serio0/input/input1
[ 0.958568] rtc_cmos 00:05: RTC can wake from S4
[ 0.960040] rtc_cmos 00:05: registered as rtc0
[ 0.961605] rtc_cmos 00:05: alarms up to one day, y3k, 242 bytes nvram, hpet irqs
[ 0.964845] device-mapper: ioctl: 4.50.0-ioctl (2025-04-28) initialised: dm-devel@lists.linux.dev
[ 0.966552] amd_pstate: the _CPC object is not present in SBIOS or ACPI disabled
[ 0.967789] hid: raw HID events driver (C) Jiri Kosina
[ 0.968993] usbcore: registered new interface driver usbhid
[ 0.969824] usbhid: USB HID core driver
```

---

```
Breakpoint 1, use_after_free () at drivers/uaf/uaf.c:7
7
(gdb) l
1
2     #include <linux/init.h>
3     #include <linux/slab.h>
4     #include <linux/string.h>
5
6     static int __init use_after_free(void)
7     {
8         char *leak = kzalloc(100, GFP_KERNEL);
9         if (!leak)
10         {
11             printk(KERN_ERR "Memory allocation failed\n");
12             goto out;
13         }
14         strcpy(leak, "This is a use-after-free example");
15         kfree(leak);
16         // Using the pointer after freeing it, which is undefined behavior
17         printk(KERN_INFO "Using freed memory: %s\n", leak); // This is dangerous!
18     out:
19         return 0;
20     }
21
(gdb) a[]
```

```
[ 0.768208] scsi 1:0:0:0: CD-ROM QEMU QEMU DVD-ROM 2.5+ PQ: 0 ANSI: 5
[ 0.769487] sd 0:0:0:0: [sda] Preferred minimum I/O size 512 bytes
[ 0.781207] sd 0:0:0:0: [sda] Attached SCSI disk
[ 0.795243] sr 1:0:0:0: [sr0] scsi3-mmc drive: 4x/4x cd/rw xa/form2 tray
[ 0.796341] cdrom: Uniform CD-ROM driver Revision: 3.20
[ 0.814677] sr 1:0:0:0: Attached scsi generic sg1 type 5
[ 0.946200] e1000 0000:00:03.0 eth0: (PCI:33MHz:32-bit) 52:54:00:12:34:56
[ 0.947302] e1000 0000:00:03.0 eth0: Intel(R) PRO/1000 Network Connection
[ 0.948390] e1000e: Intel(R) PRO/1000 Network Driver
[ 0.949128] e1000e: Copyright(c) 1999 - 2015 Intel Corporation.
[ 0.950056] sky2: driver version 1.30
[ 0.950981] usbcore: registered new interface driver usblp
[ 0.951846] usbcore: registered new interface driver usb-storage
[ 0.952800] i8042: PNP: PS/2 Controller [PNP0303:KBD,PNP0f13:MOU] at 0x60,0x64 irq 1,12
[ 0.954892] serio: i8042 KBD port at 0x60,0x64 irq 1
[ 0.955658] serio: i8042 AUX port at 0x60,0x64 irq 12
[ 0.957136] input: AT Translated Set 2 keyboard as /devices/platform/i8042/serio0/input/input1
[ 0.958568] rtc_cmos 00:05: RTC can wake from S4
[ 0.960040] rtc_cmos 00:05: registered as rtc0
[ 0.961605] rtc_cmos 00:05: alarms up to one day, y3k, 242 bytes nvram, hpet irqs
[ 0.964845] device-mapper: ioctl: 4.50.0-ioctl (2025-04-28) initialised: dm-devel@lists.linux.dev
[ 0.966552] amd_pstate: the _CPC object is not present in SBBIOS or ACPI disabled
[ 0.967789] hid: raw HID events driver (C) Jiri Kosina
[ 0.968993] usbcore: registered new interface driver usbhid
[ 0.969824] usbhid: USB HID core driver
```

---

```
8         char *leak = kzalloc(100, GFP_KERNEL);
9         if (!leak)
10         {
11             printk(KERN_ERR "Memory allocation failed\n");
(gdb)
12             goto out;
13         }
14         strcpy(leak, "This is a use-after-free example");
15         kfree(leak);
16         // Using the pointer after freeing it, which is undefined behavior
17         printk(KERN_INFO "Using freed memory: %s\n", leak); // This is dangerous!
18     out:
19         return 0;
20     }
21
(gdb) n
8         char *leak = kzalloc(100, GFP_KERNEL);
(gdb)
9         if (!leak)
(gdb)
14         strcpy(leak, "This is a use-after-free example");
(gdb)
15         kfree(leak);
(gdb)
17         printk(KERN_INFO "Using freed memory: %s\n", leak); // This is dangerous!
(gdb) 
```

```
[ 0.969824] usbhid: USB HID core driver
[ 0.972451] =====
[ 0.972909] BUG: KASAN: slab-use-after-free in string+0x476/0x570
[ 0.972909] Read of size 1 at addr ffff8880080b1300 by task swapper/0/1
[ 0.972909]
[ 0.972909] CPU: 0 UID: 0 PID: 1 Comm: swapper/0 Not tainted 6.17.7 #36 PREEMPT(voluntary)
[ 0.972909] Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS 1.17.0-5.fc42 04/01/2014
[ 0.972909] Call Trace:
[ 0.972909] <TASK>
[ 0.972909] dump_stack_lvl+0x4d/0x70
[ 0.972909] print_report+0x170/0x4f3
[ 0.972909] ? __pxf__raw_spin_lock_irqsave+0x10/0x10
[ 0.972909] kasan_report+0xda/0x110
[ 0.972909] ? string+0x476/0x570
[ 0.972909] ? string+0x476/0x570
[ 0.972909] string+0x476/0x570
[ 0.972909] ? unwind_next_frame+0x53a/0x23f0
[ 0.972909] ? arch_stack_walk+0x74/0x100
[ 0.972909] ? __pxf_string+0x10/0x10
[ 0.972909] ? __pxf_format_decode+0x10/0x10
[ 0.972909] ? unwind_next_frame+0x53a/0x23f0
[ 0.972909] vsnprintf+0x2f0/0x13f0
[ 0.972909] ? write_profile+0x90/0xf0
[ 0.972909] ? __pxf_vsnprintf+0x10/0x10
[ 0.972909] ? __raw_spin_trylock+0xaf/0x120
[ 0.972909] ? __pxf__raw_spin_trylock+0x10/0x10


---


10      {
11          printk(KERN_ERR "Memory allocation failed\n");
(gdb)         goto out;
12      }
13  }
14  strcpy(leak, "This is a use-after-free example");
15  kfree(leak);
16  // Using the pointer after freeing it, which is undefined behavior
17  printk(KERN_INFO "Using freed memory: %s\n", leak); // This is dangerous!
18  out:
19      return 0;
20  }
21
(gdb) n
8      char *leak = kzalloc(100, GFP_KERNEL);
(gdb)
9      if (!leak)
(gdb)
14      strcpy(leak, "This is a use-after-free example");
(gdb)
15      kfree(leak);
(gdb)
17      printk(KERN_INFO "Using freed memory: %s\n", leak); // This is dangerous!
(gdb) n
19      return 0;
(gdb) █
```

```
[ 0.972909] ? __pxf__raw_spin_trylock+0x10/0x10
[ 0.972909] vprintk_store+0x2c8/0x880
[ 0.972909] ? stack_depot_save_flags+0x4b5/0x8f0
[ 0.972909] ? __pxf_vprintk_store+0x10/0x10
[ 0.972909] ? kasan_save_stack+0x3f/0x50
[ 0.972909] ? __kasan_slab_free+0x63/0x70
[ 0.972909] ? kfree+0xf2/0x3b0
[ 0.972909] ? use_after_free+0x73/0x90
[ 0.972909] ? do_one_initcall+0x9d/0x310
[ 0.972909] ? kernel_init_freeable+0x458/0x4d0
[ 0.972909] vprintk_emit+0x10b/0x440
[ 0.972909] ? __pxf_vprintk_emit+0x10/0x10
[ 0.972909] ? __pxf_use_after_free+0x10/0x10
[ 0.972909] _printk+0xc2/0x100
[ 0.972909] ? __pxf__printk+0x10/0x10
[ 0.972909] ? ret_from_fork_asm+0x1a/0x30
[ 0.972909] ? ret_from_fork+0x10d/0x1b0
[ 0.972909] use_after_free+0x82/0x90 ←
[ 0.972909] do_one_initcall+0x9d/0x310
[ 0.972909] ? __pxf_do_one_initcall+0x10/0x10
[ 0.972909] ? __pxf_parse_args+0x10/0x10
[ 0.972909] ? __pxf_kasan_poison+0x10/0x10
[ 0.972909] kernel_init_freeable+0x458/0x4d0
[ 0.972909] ? __pxf_kernel_init+0x10/0x10
[ 0.972909] kernel_init+0x19/0x150
[ 0.972909] ? __pxf_kernel_init+0x10/0x10
```

```
[ 0.979210] ? kfree+0xf2/0x3b0
[ 0.979210] ? use_after_free+0x73/0x90
[ 0.979210] ? do_one_initcall+0x9d/0x310
[ 0.979210] ? kernel_init_freeable+0x458/0x4d0
[ 0.979210] vprintk_emit+0x10b/0x440
[ 0.979210] ? __pxf_vprintk_emit+0x10/0x10
[ 0.979210] ? __pxf_use_after_free+0x10/0x10
[ 0.979210] _printk+0xc2/0x100
[ 0.979210] ? __pxf__printk+0x10/0x10
[ 0.979210] ? ret_from_fork_asm+0x1a/0x30
[ 0.979210] ? ret_from_fork+0x10d/0x1b0
[ 0.979210] use_after_free+0x82/0x90
[ 0.979210] do_one_initcall+0x9d/0x310
[ 0.979210] ? __pxf_do_one_initcall+0x10/0x10
[ 0.979210] kernel_init_freeable+0x458/0x4d0
[ 0.979210] ? __pxf_kernel_init+0x10/0x10
[ 0.979210] kernel_init+0x19/0x150
[ 0.979210] ? __pxf_kernel_init+0x10/0x10
[ 0.979210] ? __pxf_kernel_init+0x10/0x10
[ 0.979210] ret_from_fork+0x10d/0x1b0
[ 0.979210] ? __pxf_kernel_init+0x10/0x10
[ 0.979210] ret_from_fork_asm+0x1a/0x30
[ 0.979210] </TASK>
[ 0.979210] Kernel Offset: disabled
[ 0.979210] ---[ end Kernel panic - not syncing: kasan.fault=panic set ... ]---
```

```
#0 delay_tsc (cycles=3892716) at arch/x86/lib/delay.c:79
#1 0xffffffff8120ce7b in udelay (usec=1000) at ./include/asm-generic/delay.h:62
#2 vpanic (fmt=fmt@entry=0xffffffff84fd1b67 "kasan.fault=panic set ...\\n", args=args@entry=0xfffff888006e2f798) at kernel/panic.c:555
#3 0xffffffff8120cf09 in panic (fmt=fmt@entry=0xffffffff84fd1b67 "kasan.fault=panic set ...\\n") at kernel/panic.c:566
#4 0xffffffff8122a348 in end_report (addr=<optimized out>, flags=0xfffff888006e2f868, is_write=false) at mm/kasan/report.c:233
#5 end_report (flags=0xfffff888006e2f868, addr=<optimized out>, is_write=<optimized out>) at mm/kasan/report.c:219
#6 0xffffffff81abbd9b in kasan_report (addr=addr@entry=0xfffff8880080b0300, size=size@entry=1, is_write=is_write@entry=false, ip=18446744071634551814) at mm/kasan/report.c:597
#7 0xffffffff81ab8a24 in __asan_report_load1_noabort (addr=addr@entry=0xfffff8880080b0300) at mm/kasan/report_generic.c:378
#8 0xffffffff84520406 in string_nocheck (buf=0xfffff888006e2fc0e "", end=<optimized out>, s=0xfffff8880080b0301 "", spec=...) at lib/vsprintf.c:653
#9 string (buf=<optimized out>, buf@entry=0xfffff888006e2fc0e "", end=end@entry=0xfffff888006e2fc00 "", s=<optimized out>, spec=...) at lib/vsprintf.c:735
#10 0xffffffff8452a6e0 in vsnprintf (buf=buf@entry=0xfffff888006e2fbff "\0016Using ", size=<optimized out>, size@entry=8, fmt_str=fmt_str@entry=0xffffffff84b058a0 "\0016Using freed memory: %s\\n",
    args=args@entry=0xfffff888006e2fb8) at lib/vsprintf.c:2926
#11 0xffffffff815f2188 in vprintk_store (facility=facility@entry=0, level=level@entry=-1, dev_info=dev_info@entry=0x0, fmt=fmt@entry=0xffffffff84b058a0 "\0016Using freed memory: %s\\n",
    args=args@entry=0xfffff888006e2fd38) at kernel/printk/printk.c:2279
#12 0xffffffff815f502b in vprintk_emit (facility=facility@entry=0, level=level@entry=-1, dev_info=dev_info@entry=0x0, fmt=fmt@entry=0xffffffff84b058a0 "\0016Using freed memory: %s\\n",
    args=args@entry=0xfffff888006e2fd38) at kernel/printk/printk.c:2426
#13 0xffffffff815f5388 in vprintk_default (fmt=fmt@entry=0xffffffff84b058a0 "\0016Using freed memory: %s\\n", args=args@entry=0xfffff888006e2fd38) at kernel/printk/printk.c:2465
#14 0xffffffff815f6c49 in vprintk (fmt=fmt@entry=0xffffffff84b058a0 "\0016Using freed memory: %s\\n", args=args@entry=0xfffff888006e2fd38) at kernel/printk/printk_safe.c:82
#15 0xffffffff812174b2 in _printk (fmt=fmt@entry=0xffffffff84b058a0 "\0016Using freed memory: %s\\n") at kernel/printk/printk.c:2475
#16 0xffffffff8630be12 in use_after_free () at drivers/uaf/uaf.c:17 ←
#17 0xffffffff8131887d in do_one_initcall (fn=0xffffffff8630bd90 <use_after_free>) at init/main.c:1281
#18 0xffffffff8624dc98 in do_initcall_level (level=6, command_line=0xfffff888006da8200 "console") at init/main.c:1343
#19 do_initcalls () at init/main.c:1359
#20 do_basic_setup () at init/main.c:1378
#21 kernel_init_freeable () at init/main.c:1591
--Type <RET> for more, q to quit, c to continue without paging--
```

pero, usemos KASAN...

el reporte de KASAN tiene:

call trace (bt),

alocaciones de memoria,

liberaciones de memoria.

en nuestro caso...

alocamos en  
**use\_after\_free+0x41/0x90**  
(slab.h)

liberamos en  
use\_after\_free+0x73/0x90

y la bt es:  
use\_after\_free+0x82/0x90

usemos gdb para encontrar...

el código ofensor.

la bt:

```
[ 0.979210] ? __pfx_string+0x10/0x10
[ 0.979210] ? __pfx_format_decode+0x10/0x10
[ 0.979210] ? unwind_next_frame+0x53a/0x23f0
[ 0.979210] vsnprintf+0x2f0/0x13f0
[ 0.979210] ? write_profile+0x90/0xf0
[ 0.979210] ? __pfx_vsnprintf+0x10/0x10
[ 0.979210] ? _raw_spin_trylock+0xaf/0x120
[ 0.979210] ? __pfx__raw_spin_trylock+0x10/0x10
[ 0.979210] vprintk_store+0x2c8/0x880
[ 0.979210] ? stack_depot_save_flags+0x4b5/0x8f0
[ 0.979210] ? __pfx_vprintk_store+0x10/0x10
[ 0.979210] ? kasan_save_stack+0x3f/0x50
[ 0.979210] ? __kasan_slab_free+0x63/0x70
[ 0.979210] ? kfree+0xf2/0x3b0
[ 0.979210] ? use_after_free+0x73/0x90
[ 0.979210] ? do_one_initcall+0x9d/0x310
[ 0.979210] ? kernel_init_freeable+0x458/0x4d0
[ 0.979210] vprintk_emit+0x10b/0x440
[ 0.979210] ? __pfx_vprintk_emit+0x10/0x10
[ 0.979210] ? __pfx_use_after_free+0x10/0x10
[ 0.979210] _printk+0xc2/0x100
[ 0.979210] ? __pfx__printk+0x10/0x10
[ 0.979210] ? ret_from_fork_asm+0x1a/0x30
[ 0.979210] ? ret_from_fork+0x10d/0x1b0
[ 0.979210] use_after_free+0x82/0x90
[ 0.979210] do_one_initcall+0x9d/0x310
```

---

```
(gdb) list *(use_after_free+0x82/0x90)
0xffffffff8630bd90 is in use_after_free (drivers/uaf/uaf.c:7).
2     #include <linux/init.h>
3     #include <linux/slab.h>
4     #include <linux/string.h>
5
6     static int __init use_after_free(void)
7     {
8         char *leak = kzalloc(100, GFP_KERNEL);
9         if (!leak)
10         {
11             printk(KERN_ERR "Memory allocation failed\n");
(gdb) 
```

allocación:

```
[ 0.979210] ? __pxf_kernel_init+0x10/0x10
[ 0.979210] ? __pxf_kernel_init+0x10/0x10
[ 0.979210] ret_from_fork+0x10d/0x1b0
[ 0.979210] ? __pxf_kernel_init+0x10/0x10
[ 0.979210] ret_from_fork_asm+0x1a/0x30
[ 0.979210] </TASK>
[ 0.979210]
[ 0.979210] Allocated by task 1:
[ 0.979210] kasan_save_stack+0x30/0x50
[ 0.979210] kasan_save_track+0x14/0x30
[ 0.979210] __kasan_kmalloc+0x9a/0xb0
[ 0.979210] use_after_free+0x41/0x90
[ 0.979210] do_one_initcall+0x9d/0x310
[ 0.979210] kernel_init_freeable+0x458/0x4d0
[ 0.979210] kernel_init+0x19/0x150
[ 0.979210] ret_from_fork+0x10d/0x1b0
[ 0.979210] ret_from_fork_asm+0x1a/0x30
[ 0.979210]
[ 0.979210] Freed by task 1:
[ 0.979210] kasan_save_stack+0x30/0x50
[ 0.979210] kasan_save_track+0x14/0x30
[ 0.979210] kasan_save_free_info+0x3b/0x70
[ 0.979210] __kasan_slab_free+0x63/0x70
[ 0.979210] kfree+0xf2/0x3b0
[ 0.979210] use_after_free+0x73/0x90
[ 0.979210] do_one_initcall+0x9d/0x310
(gdb) list *(use_after_free+0x90)
(gdb) list *(use_after_free+0x41/0x90)
0xffffffff8630bd90 is in use_after_free (drivers/uaf/uaf.c:7).
2     #include <linux/init.h>
3     #include <linux/slab.h>
4     #include <linux/string.h>
5
6     static int __init use_after_free(void)
7     {
8         char *leak = kzalloc(100, GFP_KERNEL);
9         if (!leak)
10         {
11             printk(KERN_ERR "Memory allocation failed\n");
(gdb) █
```

liberación:

```
[ 0.979210] </TASK>
[ 0.979210]
[ 0.979210] Allocated by task 1:
[ 0.979210] kasan_save_stack+0x30/0x50
[ 0.979210] kasan_save_track+0x14/0x30
[ 0.979210] __kasan_kmalloc+0x9a/0xb0
[ 0.979210] use_after_free+0x41/0x90
[ 0.979210] do_one_initcall+0x9d/0x310
[ 0.979210] kernel_init_freeable+0x458/0x4d0
[ 0.979210] kernel_init+0x19/0x150
[ 0.979210] ret_from_fork+0x10d/0x1b0
[ 0.979210] ret_from_fork_asm+0x1a/0x30
[ 0.979210]
[ 0.979210] Freed by task 1:
[ 0.979210] kasan_save_stack+0x30/0x50
[ 0.979210] kasan_save_track+0x14/0x30
[ 0.979210] kasan_save_free_info+0x3b/0x70
[ 0.979210] __kasan_slab_free+0x63/0x70
[ 0.979210] kfree+0xf2/0x3b0
[ 0.979210] use_after_free+0x73/0x90
[ 0.979210] do_one_initcall+0x9d/0x310
[ 0.979210] kernel_init_freeable+0x458/0x4d0
[ 0.979210] kernel_init+0x19/0x150
[ 0.979210] ret_from_fork+0x10d/0x1b0
[ 0.979210] ret_from_fork_asm+0x1a/0x30
[ 0.979210]

(gdb) list *(use_after_free+0x73)
0xffffffff8630be03 is in use_after_free (drivers/uaf/uaf.c:17).
12         goto out;
13     }
14     strcpy(leak, "This is a use-after-free example");
15     kfree(leak);
16     // Using the pointer after freeing it, which is undefined behavior
17     printk(KERN_INFO "Using freed memory: %s\n", leak); // This is dangerous!
18     out:
19     return 0;
20 }
21
```

es necesario jugar con los  
valores...

para encontrar los bugs.

tracing.

muchá depuración, ahora  
pasemos a la exploración.

ftrace: el papi de strace.

muchos conocemos strace.

pero ftrace está un paso más abajo.

the function tracer.

pasamos de solo ver  
syscalls...

a ver el código completo.

ftrace es parte de la  
arquitectura del kernel.

basado en callbacks.

trace-cmd

podríamos usar ftrace a través  
de la interfaz...

sysfs /sys/kernel/tracing.

para eso existe trace-cmd.

veamos el clásico printf.

hola amigos

[ anti@mothership ]  
11:41 AM

~/Documents/Research/linux-debugging/ftrace/code\_example \$ |

strace cuenta una historia...

```
~/Documents/Research/linux-debugging/ftrace/code_example $ strace ./main
execve("./main", ["../main"], 0x7ffe150f1400 /* 61 vars */) = 0
brk(NULL)                                = 0x3eed1000
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f2496893000
access("/etc/ld.so.preload", R_OK)        = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=145907, ...}) = 0
mmap(NULL, 145907, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f249686f000
close(3)                                 = 0
openat(AT_FDCWD, "/lib64/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\0\0>\0\1\0\0\00007\0\0\0\0\0\0"..., 832) = 832
pread64(3, "\6\0\0\0\4\0\0@0\0\0\0\0\0@0\0\0\0\0\0@0\0\0\0\0@0\0\0\0\0"..., 784, 64) = 784
fstat(3, {st_mode=S_IFREG|0755, st_size=2447520, ...}) = 0
pread64(3, "\6\0\0\0\4\0\0@0\0\0\0\0@0\0\0\0\0@0\0\0\0\0@0\0\0\0\0"..., 784, 64) = 784
mmap(NULL, 2038832, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x7f249667d000
mmap(0x7f24967ec000, 479232, PROT_READ, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x16f000) = 0x7f24967ec000
mmap(0x7f2496861000, 24576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x1e3000) = 0x7f2496861000
mmap(0x7f2496867000, 31792, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0x7f2496867000
close(3)                                 = 0
mmap(NULL, 12288, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f249667a000
arch_prctl(ARCH_SET_FS, 0x7f249667a740) = 0
set_tid_address(0x7f249667aa10)          = 2331067
set_robust_list(0x7f249667aa20, 24)     = 0
rseq(0x7f249667a680, 0x20, 0, 0x53053053) = 0
mprotect(0x7f2496861000, 16384, PROT_READ) = 0
mprotect(0x402000, 4096, PROT_READ)       = 0
mprotect(0x7f24968d1000, 8192, PROT_READ) = 0
prlimit64(0, RLIMIT_STACK, NULL, {rlim_cur=16384*1024, rlim_max=RLIM64_INFINITY}) = 0
munmap(0x7f249686f000, 145907)          = 0
fstat(1, {st_mode=S_IFCHR|0620, st_rdev=makedev(0x88, 0x2), ...}) = 0
getrandom("\xaa\x3d\x8a\xc3\xff\x17\xe3\x62", 8, GRND_NONBLOCK) = 8
brk(NULL)                                = 0x3eed1000
brk(0x3eef2000)                          = 0x3eef2000
write(1, "hola amigos\n", 12)             ←
)                                         = 12
exit_group(0)                            = ?
+++ exited with 0 +++
```

interesante...

pero limitado.

pasemos a ftrace.

```
~/Documents/Research/linux-debugging/ftrace/code_example $ sudo trace-cmd record -p function_graph -F ./main
  plugin 'function_graph'
hola amigos
CPU0 data recorded at offset=0x3ef000
  0 bytes in size (0 uncompressed)
CPU1 data recorded at offset=0x3ef000
  0 bytes in size (0 uncompressed)
CPU2 data recorded at offset=0x3ef000
  0 bytes in size (0 uncompressed)
CPU3 data recorded at offset=0x3ef000
  0 bytes in size (0 uncompressed)
CPU4 data recorded at offset=0x3ef000
  0 bytes in size (0 uncompressed)
CPU5 data recorded at offset=0x3ef000
  0 bytes in size (0 uncompressed)
CPU6 data recorded at offset=0x3ef000
  309893 bytes in size (1802240 uncompressed)
CPU7 data recorded at offset=0x43b000
  0 bytes in size (0 uncompressed)
CPU8 data recorded at offset=0x43b000
  0 bytes in size (0 uncompressed)
CPU9 data recorded at offset=0x43b000
  0 bytes in size (0 uncompressed)
CPU10 data recorded at offset=0x43b000
  0 bytes in size (0 uncompressed)
CPU11 data recorded at offset=0x43b000
  0 bytes in size (0 uncompressed)
```

hmm . . .

```
~/Documents/Research/linux-debugging/ftrace/code_example $ sudo trace-cmd report | wc -l  
29188
```

¡¿29000 funciones!?

podemos decir que el output es  
un tanto más grande.

cpus=12

main-2331557 [006] ...2. 227930.190666: funcgraph_entry:	1.222 us	mutex_unlock(); (ret=0xfffff8e4e5461b080)
main-2331557 [006] ...2. 227930.190668: funcgraph_entry:		__f_unlock_pos() {
main-2331557 [006] ...2. 227930.190668: funcgraph_exit:	0.080 us	mutex_unlock(); (ret=0xfffff8e4e5461b080)
main-2331557 [006] ...1. 227930.190668: funcgraph_exit:	0.250 us	} (ret=0xfffff8e4e5461b080)
main-2331557 [006] ...2. 227930.190668: funcgraph_entry:	0.061 us	syscall_exit_to_user_mode_prepare(); (ret=0x1)
main-2331557 [006] ...2. 227930.190669: funcgraph_entry:		__x64_sys_execve()
main-2331557 [006] ...2. 227930.190669: funcgraph_entry:		getname_flags()
main-2331557 [006] ...2. 227930.190669: funcgraph_entry:	0.100 us	getname_flags.part.0() {
main-2331557 [006] ...2. 227930.190669: funcgraph_entry:		kmem_cache_alloc_noprof(); (ret=0xfffff8e4c826a7000)
main-2331557 [006] ...2. 227930.190669: funcgraph_entry:		__check_object_size()
main-2331557 [006] ...2. 227930.190669: funcgraph_entry:	0.060 us	__check_object_size.part.0() {
main-2331557 [006] ...2. 227930.190670: funcgraph_entry:	0.070 us	check_stack_object(); (ret=0x0)
main-2331557 [006] ...2. 227930.190670: funcgraph_entry:	0.070 us	is_vmalloc_addr(); (ret=0x0)
main-2331557 [006] ...2. 227930.190670: funcgraph_entry:	0.070 us	__virt_addr_valid(); (ret=0x1)
main-2331557 [006] ...2. 227930.190670: funcgraph_entry:	0.050 us	__check_heap_object(); (ret=0xfe0)
main-2331557 [006] ...1. 227930.190670: funcgraph_exit:	0.651 us	} (ret=0xfffff8e5305200000)
main-2331557 [006] ...1. 227930.190670: funcgraph_exit:	0.792 us	} (ret=0xfffff8e5305200000)
main-2331557 [006] ...1. 227930.190670: funcgraph_exit:	1.172 us	} (ret=0xfffff8e4c826a7000)
main-2331557 [006] ...1. 227930.190670: funcgraph_exit:	1.302 us	} (ret=0xfffff8e4c826a7000)
main-2331557 [006] ...2. 227930.190671: funcgraph_entry:		do_execveat_common.isra.0() {
main-2331557 [006] ...2. 227930.190671: funcgraph_entry:		alloc_bprm()
main-2331557 [006] ...2. 227930.190671: funcgraph_entry:		do_open_execat()
main-2331557 [006] ...2. 227930.190671: funcgraph_entry:		do_filp_open()
main-2331557 [006] ...2. 227930.190671: funcgraph_entry:		path_openat()
main-2331557 [006] ...2. 227930.190671: funcgraph_entry:		alloc_empty_file()
main-2331557 [006] ...2. 227930.190671: funcgraph_entry:		kmem_cache_alloc_noprof()
main-2331557 [006] ...2. 227930.190671: funcgraph_entry:	0.080 us	__memcg_slab_post_alloc_hook()
main-2331557 [006] ...2. 227930.190671: funcgraph_entry:		obj_cgroup_charge()
main-2331557 [006] ...2. 227930.190671: funcgraph_entry:	0.230 us	consume_obj_stock(); (ret=0x1)
main-2331557 [006] ...1. 227930.190672: funcgraph_exit:	0.230 us	} (ret=0x0)
main-2331557 [006] ...2. 227930.190672: funcgraph_entry:	0.070 us	__rcu_read_lock(); (ret=0x1)
main-2331557 [006] ...2. 227930.190672: funcgraph_entry:	0.060 us	__rcu_read_unlock(); (ret=0x0)
main-2331557 [006] ...2. 227930.190672: funcgraph_entry:	0.090 us	mod_objcg_state(); (ret=0xfffff8e5bb7b70000)
main-2331557 [006] ...1. 227930.190672: funcgraph_exit:	0.881 us	} (ret=0x1)
main-2331557 [006] ...1. 227930.190672: funcgraph_exit:	1.032 us	} (ret=0xfffff8e5373bde6c0)
main-2331557 [006] ...2. 227930.190672: funcgraph_entry:		init_file()
main-2331557 [006] ...2. 227930.190672: funcgraph_entry:		security_file_alloc()
main-2331557 [006] ...2. 227930.190672: funcgraph_entry:	0.080 us	kmem_cache_alloc_noprof(); (ret=0xfffff8e5064c32260)
main-2331557 [006] ...2. 227930.190673: funcgraph_entry:	0.060 us	selinux_file_alloc_security(); (ret=0x0)
main-2331557 [006] ...2. 227930.190673: funcgraph_entry:	0.060 us	bpf_lsm_file_alloc_security(); (ret=0x0)
main-2331557 [006] ...2. 227930.190673: funcgraph_entry:	0.070 us	hook_file_alloc_security(); (ret=0x0)
main-2331557 [006] ...1. 227930.190673: funcgraph_exit:	0.641 us	} (ret=0x0)
main-2331557 [006] ...2. 227930.190673: funcgraph_entry:	0.070 us	__mutex_init(); (ret=0xfffff8e5373bde720)
main-2331557 [006] ...1. 227930.190673: funcgraph_exit:	0.942 us	} (ret=0x0)
main-2331557 [006] ...1. 227930.190673: funcgraph_exit:	2.184 us	} (ret=0xfffff8e5373bde6c0)
main-2331557 [006] ...2. 227930.190673: funcgraph_entry:		path_init()
main-2331557 [006] ...2. 227930.190673: funcgraph_entry:		__rcu_read_lock(); (ret=0x1)
main-2331557 [006] ...1. 227930.190674: funcgraph_exit:	0.070 us	} (ret=0xfffff8e4c826a7020)
main-2331557 [006] ...2. 227930.190674: funcgraph_entry:	0.250 us	link_path_walk.part.0.constprop.0() {
main-2331557 [006] ...2. 227930.190674: funcgraph_entry:		inode_permission()

no veremos el output  
completo...

pero nos sirve para hacernos  
una idea.

del poder de ftrace.

trace-cmd; profundizando .

los principales comandos que usaremos . . .

son trace-cmd record...

y trace-cmd report.

por lo general, utilizaremos  
flags en trace-cmd record.

man trace-cmd-record  
nos enseña las flags.

pero veamos algunas.

-F programa  
filtra por llamadas hechas por  
solamente lo que vayamos a  
ejecutar.

-p plugin  
seleccionamos el plugin de  
visualización (generalmente  
function\_graph).

-n funcion  
nos permite "ignorar" ciertas  
trazas.

-g funcion  
nos permite enfocar el tracing a ciertas funciones.

-o archivo  
guardar el trace.dat en un  
archivo 'archivo'.

--max-graph-depth n  
profundidad máxima de funciones.

ahora . . .

¡combinemos!

```
sudo trace-cmd record \
-p function_graph \
-g __x64_sys_write \
-n rw_verify_area \
-F ./main
```

```
~/Documents/Research/linux-debugging/ftrace/code_example $ trace-cmd report | wc -l  
273
```

273 es un valor más entendible

la idea no es leer 29k funciones

explorar y minimizar el alcance.

kernelshark

en vez de usar trace-cmd report  
para observar la traza...

podemos utilizar kernelshark.

como el nombre indica...

es como wireshark.

pero para trazas del kernel.

cpus=12

main-2361126 [007] ...2. 239669.404589: funcgraph_entry:	do_output_char()
main-2361126 [007] ...2. 239669.404590: funcgraph_entry:	pty_write()
main-2361126 [007] ...2. 239669.404590: funcgraph_entry:	tty_insert_flip_string_and_push_buffer()
main-2361126 [007] ...2. 239669.404590: funcgraph_entry:	__raw_spin_lock_irqsave(); (ret=0x292)
main-2361126 [007] d..3. 239669.404591: funcgraph_entry:	__tty_insert_flip_string_flags()
main-2361126 [007] d..3. 239669.404591: funcgraph_entry:	__tty_buffer_request_room(); (ret=0x2)
main-2361126 [007] d..2. 239669.404591: funcgraph_exit:	} (ret=0x2)
main-2361126 [007] d..3. 239669.404591: funcgraph_entry:	__raw_spin_unlock_irqrestore(); (ret=0xfffff8e4cdb3f3800)
main-2361126 [007] ...2. 239669.404591: funcgraph_entry:	queue_work_on()
main-2361126 [007] d..2. 239669.404591: funcgraph_entry:	clear_pending_if_disabled(); (ret=0x0)
main-2361126 [007] d..2. 239669.404591: funcgraph_entry:	__queue_work()
main-2361126 [007] d..2. 239669.404591: funcgraph_entry:	__rcu_read_lock(); (ret=0x1)
main-2361126 [007] d..2. 239669.404591: funcgraph_entry:	wq_select_unbound_cpu(); (ret=0x7)
main-2361126 [007] d..2. 239669.404591: funcgraph_entry:	__raw_spin_lock(); (ret=0x0)
main-2361126 [007] d..3. 239669.404591: funcgraph_entry:	pwq_tryinc_nr_active(); (ret=0x1)
main-2361126 [007] d..3. 239669.404591: funcgraph_entry:	insert_work(); (ret=0x3)
main-2361126 [007] d..3. 239669.404592: funcgraph_entry:	kick_pool()
main-2361126 [007] d..3. 239669.404592: funcgraph_entry:	wake_up_process()
main-2361126 [007] d..3. 239669.404592: funcgraph_entry:	try_to_wake_up()
main-2361126 [007] d..4. 239669.404592: funcgraph_entry:	__raw_spin_lock_irqsave(); (ret=0x82)
main-2361126 [007] d..6. 239669.404592: funcgraph_entry:	__traceiter_sched_waking()
main-2361126 [007] d..6. 239669.404592: funcgraph_entry:	__raw_spin_lock_irqsave(); (ret=0x82)
main-2361126 [007] d..7. 239669.404592: funcgraph_entry:	__raw_spin_unlock_irqrestore(); (ret=0xfffff8e4c974b9000)
main-2361126 [007] d..5. 239669.404592: funcgraph_exit:	} (ret=0x0)
main-2361126 [007] d..5. 239669.404592: funcgraph_entry:	select_task_rq_fair()
main-2361126 [007] d..5. 239669.404592: funcgraph_entry:	__rcu_read_lock(); (ret=0x2)
main-2361126 [007] d..5. 239669.404593: funcgraph_entry:	wake_affine()
main-2361126 [007] d..5. 239669.404593: funcgraph_entry:	available_idle_cpu(); (ret=0x0)
main-2361126 [007] d..5. 239669.404593: funcgraph_entry:	available_idle_cpu(); (ret=0x0)
main-2361126 [007] d..5. 239669.404593: funcgraph_entry:	task_h_load(); (ret=0x0)
main-2361126 [007] d..4. 239669.404593: funcgraph_exit:	} (ret=0xb)
main-2361126 [007] d..5. 239669.404593: funcgraph_entry:	select_idle_sibling()
main-2361126 [007] d..5. 239669.404593: funcgraph_entry:	available_idle_cpu(); (ret=0x0)
main-2361126 [007] d..5. 239669.404593: funcgraph_entry:	cpus_share_cache(); (ret=0xfffff8e5bb7df0001)
main-2361126 [007] d..5. 239669.404593: funcgraph_entry:	available_idle_cpu(); (ret=0x0)
main-2361126 [007] d..5. 239669.404594: funcgraph_entry:	select_idle_cpu()
main-2361126 [007] d..5. 239669.404594: funcgraph_entry:	select_idle_core.isra.0()
main-2361126 [007] d..5. 239669.404594: funcgraph_entry:	available_idle_cpu(); (ret=0x1)
main-2361126 [007] d..5. 239669.404594: funcgraph_entry:	available_idle_cpu(); (ret=0x1)
main-2361126 [007] d..4. 239669.404594: funcgraph_exit:	} (ret=0x0)
main-2361126 [007] d..4. 239669.404594: funcgraph_exit:	} (ret=0x0)
main-2361126 [007] d..4. 239669.404594: funcgraph_exit:	__rcu_read_unlock(); (ret=0x0)
main-2361126 [007] d..5. 239669.404594: funcgraph_entry:	} (ret=0x0)
main-2361126 [007] d..4. 239669.404594: funcgraph_exit:	kthread_is_per_cpu(); (ret=0x0)
main-2361126 [007] d..5. 239669.404594: funcgraph_entry:	set_task_cpu()
main-2361126 [007] d..5. 239669.404595: funcgraph_entry:	migrate_task_rq_fair() {
main-2361126 [007] d..5. 239669.404595: funcgraph_entry:	}

trace-cmd report, feo!

Pointer: [ ]

&lt; + - &gt; ++ -- Marker A

Marker B

A,B Delta: [ ]

239669.404581

239669.404596

239669.404610

CPU 7

Search: Column # contains Next Prev  Graph follows

#	CPU	Time Stamp	Task	PID	Latency	Event	Info
0	7	239669.404589	main	2361126	...2.	ftrace/funcgraph_entry	do_output_char() {
1	7	239669.404590	main	2361126	...2.	ftrace/funcgraph_entry	pty_write() {
2	7	239669.404590	main	2361126	...2.	ftrace/funcgraph_entry	tty_insert_flip_string_and_push_buffer() {
3	7	239669.404590	main	2361126	...2.	ftrace/funcgraph_exit	0.070 us   _raw_spin_lock_irqsave(); (ret=0x292)
4	7	239669.404590	main	2361126	d..2.	ftrace/funcgraph_exit	0.070 us   } (ret=0x292)
5	7	239669.404591	main	2361126	d..3.	ftrace/funcgraph_entry	__tty_insert_flip_string_flags() {
6	7	239669.404591	main	2361126	d..3.	ftrace/funcgraph_entry	0.040 us   __tty_buffer_request_room(); (ret=0x2)
7	7	239669.404591	main	2361126	d..2.	ftrace/funcgraph_exit	0.040 us   } (ret=0x2)
8	7	239669.404591	main	2361126	d..2.	ftrace/funcgraph_exit	0.160 us   } (ret=0x2)
9	7	239669.404591	main	2361126	d..3.	ftrace/funcgraph_entry	0.050 us   _raw_spin_unlock_irqrestore(); (ret=0xffff8e4cdb3f3800)
10	7	239669.404591	main	2361126	...1.	ftrace/funcgraph_exit	0.050 us   } (ret=0xffff8e4cdb3f3800)
11	7	239669.404591	main	2361126	...2.	ftrace/funcgraph_entry	queue_work_on() {
12	7	239669.404591	main	2361126	d..2.	ftrace/funcgraph_entry	0.060 us   clear_pending_if_disabled(); (ret=0x0)
13	7	239669.404591	main	2361126	d..1.	ftrace/funcgraph_exit	0.060 us   } (ret=0x0)
14	7	239669.404591	main	2361126	d..2.	ftrace/funcgraph_entry	__queue_work() {
15	7	239669.404591	main	2361126	d..2.	ftrace/funcgraph_entry	0.040 us   __rcu_read_lock(); (ret=0x1)
16	7	239669.404591	main	2361126	d..1.	ftrace/funcgraph_exit	0.040 us   } (ret=0x1)
17	7	239669.404591	main	2361126	d..2.	ftrace/funcgraph_entry	0.050 us   wq_select_unbound_cpu(); (ret=0x7)
18	7	239669.404591	main	2361126	d..1.	ftrace/funcgraph_exit	0.050 us   } (ret=0x7)
19	7	239669.404591	main	2361126	d..2.	ftrace/funcgraph_entry	0.050 us   _raw_spin_lock(); (ret=0x0)
20	7	239669.404591	main	2361126	d..2.	ftrace/funcgraph_exit	0.050 us   } (ret=0x0)
21	7	239669.404591	main	2361126	d..3.	ftrace/funcgraph_entry	0.051 us   pwq_tryinc_nr_active(); (ret=0x1)
22	7	239669.404591	main	2361126	d..2.	ftrace/funcgraph_exit	0.051 us   } (ret=0x1)
23	7	239669.404591	main	2361126	d..3..	ftrace/funcgraph_entry	0.040 us   insert_work(); (ret=0x3)
24	7	239669.404592	main	2361126	d..2..	ftrace/funcgraph_exit	0.040 us   } (ret=0x3)
25	7	239669.404592	main	2361126	d..3..	ftrace/funcgraph_entry	kick_pool() {
26	7	239669.404592	main	2361126	d..3..	ftrace/funcgraph_entry	wake_up_process() {
27	7	239669.404592	main	2361126	d..3..	ftrace/funcgraph_entry	try_to_wake_up() {
28	7	239669.404592	main	2361126	d..4..	ftrace/funcgraph_entry	0.040 us   _raw_spin_lock_irqsave(); (ret=0x82)
29	7	239669.404592	main	2361126	d..4..	ftrace/funcgraph_exit	0.040 us   } (ret=0x82)
30	7	239669.404592	main	2361126	d..6..	ftrace/funcgraph_entry	__traceiteiter_sched_waking() {
31	7	239669.404592	main	2361126	d..6..	ftrace/funcgraph_entry	0.040 us   _raw_spin_lock_irqsave(); (ret=0x82)
32	7	239669.404592	main	2361126	d..6..	ftrace/funcgraph_exit	0.040 us   } (ret=0x82)
33	7	239669.404592	main	2361126	d..7..	ftrace/funcgraph_entry	0.050 us   _raw_spin_unlock_irqrestore(); (ret=0xffff8e4c974b9000)
34	7	239669.404592	main	2361126	d..5..	ftrace/funcgraph_exit	0.050 us   } (ret=0xffff8e4c974b9000)
35	7	239669.404592	main	2361126	d..5..	ftrace/funcgraph_exit	0.310 us   } (ret=0x0)

kernelshark! bonito!

outro

hoy hemos pasado por muchas cosas.

desde compilación...

flags de depuración...

uso de GDB+KGDB . . .

depuramos un UAF . . .

vimos ftrace...

trace-cmd . . .

y kernelshark.

adios.  
un besito a todos.

samuel@securejump.cl  
4nt1@resacachile.cl  
@4nt1@lile.cl  
@clp\_c tg

un besito a mi esposa tambien